
December 3rd, 2024 at 15:24 UTC 2

Copyleaks Analysis Report

Al Detection Report

New Scans 3:24 PM

Plagiarism Detection N/A Plagiarism Types Identical Minor Changes Paraphrased Excluded Omitted Words



Total Pages:

Total Words:

285







Al Content



	Text Coverage	Words			
Al Text	100%	285			
Low FrequencyMedium FrequencyHigh Frequency		40 9 4			
			Human Text	0%	0
			Excluded		
Omitted Words		0			

About Al Detection

Our AI Detector is the only enterprise-level solution that can verify if the content was written by a human or generated by AI, including source code and text that has been plagiarized or modified. <u>Learn more</u>

Al Text

Human Text

A body of text that has been generated or altered by AI technology. Learn more

Any text that has been fully written by a human and has not been altered or generated by Al. <u>Learn more</u>

Copyleaks AI Detector Effectiveness

Credible data at scale, coupled with machine learning and widespread adoption, allows us to continually refine and improve our ability to understand complex text patterns, resulting in over 99% accuracy—far higher than any other AI detector—and improving daily. <u>Learn more</u>

Ideal Text Length

The higher the character count, the easier for our technology to determine irregular patterns, which results in a higher confidence rating for AI detection. Learn more

Reasons It Might Be AI When You Think It's Not

The AI Detector can detect a variety of AI-generated text, including tools that use AI technology to paraphrase content, auto-complete sentences, and more. Learn more

User AI Alert History

Historical data of how many times a user has been flagged for potentially having AI text within their content. Learn more

Al Insights

 $The number of times a phrase was found more frequently in AI vs human text is shown according to low, medium, and high frequency. \underline{Learn more}$

The frequency of a phrase in AI vs. human text.

7 x 156x

156x Existing Methods: The paper

How frequently the phrase was found in our dataset:

Al Text 5.98 / 1,000,000 Documents

Human Text 0.04 / 1,000,000 Documents

50x noise and redundant information

How frequently the phrase was found in our dataset:

Al Text 3.19 / 1,000,000 Documents

Human Text 0.06 / 1,000,000 Documents

32x noise and redundancy

How frequently the phrase was found in our dataset:

Al Text 7.3 / 1,000,000 Documents

Human Text 0.23 / 1,000,000 Documents

22x further highlights its

How frequently the phrase was found in our dataset:

Al Text 5.36 / 1,000,000 Documents

Human Text 0.24 / 1,000,000 Documents

16x demonstrates superior performance.

How frequently the phrase was found in our dataset:

Al Text 11.01 / 1,000,000 Documents

Human Text 0.71 / 1,000,000 Documents

12x and the potential limitations of

How frequently the phrase was found in our dataset:

Al Text 3.17 / 1,000,000 Documents

Human Text 0.26 / 1,000,000 Documents

8x multiple datasets and

How frequently the phrase was found in our dataset:

Al Text 12.17 / 1,000,000 Documents

Human Text 1.46 / 1,000,000 Documents

7x reveal different aspects of

How frequently the phrase was found in our dataset:

 Al Text
 7.99 / 1,000,000 Documents

 Human Text
 1.18 / 1,000,000 Documents

59x on the challenges faced during

How frequently the phrase was found in our dataset:

Al Text 3.71 / 1,000,000 Documents

Human Text 0.06 / 1,000,000 Documents

35x for larger datasets.

How frequently the phrase was found in our dataset:

Al Text 50.17 / 1,000,000 Documents

Human Text 1.43 / 1,000,000 Documents

31x robustness and generalization?

How frequently the phrase was found in our dataset:

Al Text 15.74 / 1,000,000 Documents

Human Text 0.51 / 1,000,000 Documents

17x authors provide insights into

How frequently the phrase was found in our dataset:

Al Text 2.55 / 1,000,000 Documents

Human Text 0.15 / 1,000,000 Documents

13x mathematical formulations and

How frequently the phrase was found in our dataset:

Al Text 7.99 / 1,000,000 Documents

Human Text 0.63 / 1,000,000 Documents

10x Weaknesses Lack of

How frequently the phrase was found in our dataset:

Al Text 16.45 / 1,000,000 Documents

Human Text 1.61 / 1,000,000 Documents

7x Case Studies: While

How frequently the phrase was found in our dataset:

Al Text 11.75 / 1,000,000 Documents

Human Text 1.59 / 1,000,000 Documents

Strengths

Innovative Framework: The proposed DVIB framework integrates variational information bottleneck techniques with hidden-layer perturbations and self-distillation. This approach is both novel and theoretically sound, offering a unified solution for addressing noise and redundancy issues in multimodal recommender systems.

Theoretical Justification: The authors provide robust theoretical evidence for the efficacy of DVIB. The mathematical formulations and proofs convincingly demonstrate how the approach reduces noise and redundant information while improving feature extraction quality.

Comprehensive Evaluation: The experiments span multiple datasets and baseline models, showing consistent performance improvements. The adaptability of DVIB with both constant and adaptive noise further highlights its practical utility.

Weaknesses

Lack of Real-World Case Studies: While the theoretical and experimental settings are rigorous, the absence of real-world deployment scenarios limits the practical demonstration of the proposed framework's effectiveness.

Computational Cost Analysis: Although the paper mentions no additional inference cost, the training efficiency gains need more detailed exploration, particularly concerning scalability for larger datasets.

Limited Discussion on Generalization: The adaptability of DVIB to other domains beyond multimodal recommender systems, such as general machine learning tasks, is not addressed.

Questions:

1.Integration with Existing Methods: The paper shows compatibility with other robustness-enhancing methods (e.g., AMR, MG). Could you elaborate on the challenges faced during such integrations and the potential limitations of DVIB when applied alongside complex existing pipelines? 2. Adaptive Noise Scalability: The adaptive noise mechanism demonstrates superior performance. However, how does this method scale with increasingly large multimodal datasets or higher-dimensional features? Are there any observed bottlenecks? 3. Evaluation Metrics Beyond Top-5: The experiments primarily focus on Top-5 metrics. Can the authors provide insights into the framework's performance for broader rankings or less popular recommendations, which might reveal different aspects of robustness and generalization?